

#	Area	Project
1	Pre-sales	Goal based insurance simulator using Agentic AI (voice/ text based)
2	Pre-sales	Humanoid calling to sales agent to create dossier on leads
3	Sales	Micro-market data enrichment using non-conventional signals
4	Sales	Enable sales conversations using real-time pitch assistant based on Agentic workflow
5	UW	L1 underwriting performed by Agentic AI
6	UW	Video based lie detector
7	Issuance	Issuance workflow replication using Agentic AI
8	AUM	Bionic calling
9	Claims	Agentic AI platform for evaluation and processing of claims
10	Product	Intelligent product comparison and gap analysis tool
11	Compliance	Central repository for accessing notifications/ regulations for regulators/ LI peers
12	BOL	Real-time video assistant on website

1. Goal-Based Insurance Simulator

Objective

Develop an interactive, GenAI-powered simulator that helps customers **visualize their future life goals** (such as a child's education, buying a home, or retirement). The tool should **suggest the right life insurance products** to achieve these goals, moving the experience from "selling a policy" to "planning a life."

Problem statement

Currently, most insurance websites only offer basic calculators or generic product lists. Customers often find it difficult to understand exactly how much insurance or investment they need for specific life milestones. There is a gap between a customer's real-life aspirations and the complex financial products available. We need a solution that makes financial planning conversational, visual, and highly personalized.

Core Project Scope & Deliverables

1. Automated Product Engine (The Data)

- **Dynamic Data Fetching:** The tool must fetch product details (brochures/features) directly from the official ICICI Prudential website.
- **Monthly Refresh:** Ensure the engine dynamically refreshes product information every month so recommendations are always based on the latest offerings.

2. Conversational GenAI Interface (The Input)

- Build a chat-based interface where the user provides "Life Context":
 - **Demographics:** Age, city, and family structure.
 - **Financials:** Income and risk appetite.
 - **Aspirations:** Goals like "Child's Higher Education," "Marriage," or "Retirement."

3. "What-If" Simulation Engine (The Logic)

- The tool must allow users to simulate different scenarios to see how it affects their financial needs.
 - *Example:* "What if I want to retire at 50 instead of 60?" or "What if I want to send my child to study abroad instead of locally?"
- The backend should calculate the required "corpus" (total money needed) versus current savings and identify the gap.

4. Product Mapping & Personalized Roadmap (The Output)

- **Smart Recommendation:** Match the customer's goals with the specific IPru products that fit best (e.g., ULIP for wealth, Term plans for protection, Annuities for retirement).
- **Visual Life Journey:** Instead of a text list, generate a visual dashboard (like a timeline or graph) showing the customer's "Life Journey" and how the suggested insurance plan covers each milestone.

2. Humanoid calling to sales agent to create dossier on leads

Objective

Build a voice-AI interface that allows an FSC to **"create a dossier"** on a leads/potential customer meeting via a guided audio interview. The AI must transcribe the conversation, extract specific affluence data points, and save them into a structured database.

Key Technical Challenges to Solve:

1. **Multilingual NLP:** The system must handle "Hinglish" (mix of Hindi/English) and regional dialects seamlessly.
2. **Acoustic Resilience:** Superior noise cancellation to filter out heavy background noise (traffic, wind while on a bike, or cafe chatter).
3. **Conversational Intelligence:** The bot must handle audio dropouts or mid-interview interruptions and resume gracefully.

Requirement: The "Logic-Based" Interview

AI must use conditional branching. It should only ask follow-up questions if the previous answer makes them relevant.

The 15 Prioritized Data Points & Logic Flow:

1. **Professional Profile:** What is the customer's profession and specific job role/business nature?
2. **Income (Individual):** What is their annual individual income?
3. **Family Structure:** Is the customer married? (*Condition: If Yes, proceed to Household Income & Kids*).
4. **Household Income:** Including the spouse/others, what is the total annual household income?
5. **Children:** Does the customer have children?
How old are the children? (*Condition: If Yes, proceed to Education*).
6. **Education Quality:** What kind of school (CBSE/ ICSE/ IB) or college (public/ pvt/ India/ abroad) the children are studying in? If they are transitioning from school to college - what are the higher education plans?
7. **Lifestyle Markers:** What are the customer's primary hobbies or passions (e.g., luxury collections, travel)?
8. **Travel Frequency:** How often do they travel and what is the typical spend per trip?
9. **Social Status:** Does the customer hold any premium club memberships?
10. **Residence Status:** Does the customer own or rent their home?
11. **Housing Type:** Is it a flat, villa, or independent bungalow? (*Condition: If Owned, proceed to Home Value*).
12. **Real Estate Value:** What is the approximate market value of their current residence?
13. **Vehicle Ownership:** Does the customer own a vehicle? (*Condition: If Yes, ask for Make/Model*).
14. **Asset Class:** What is the specific make and model of the car/two-wheeler?
15. **Wealth & Goals:** How many properties do they own in total and what is their primary financial goal (e.g., wealth creation vs. child's marriage)?

Deliverables:

- A functional voice-bot prototype capable of the branching logic described above (adaptive probing).
- An output file (Excel/JSON) showing the data extracted from a "noisy" test recording.
- A brief technical note on how the system handles "Hinglish" and regional accents.

3. Micro-market data enrichment

Objective

Develop a **data-driven framework to identify affluent geographies in India at a micro-market level (pincode or smaller)** using **non-conventional signals**. The goal is to build a **scalable model that ranks Indian pincodes by affluence**, enabling targeted distribution and product strategy.

Problem statement

India lacks reliable public micro-level income data. Therefore, teams must infer affluence using proxy signals. Participants must design a data pipeline + scoring model that:

- Collects publicly available signals
- Processes them geographically
- Builds an affluence scoring model
- Produces ranked pincode-level insights

Some of the non-conventional signals that may be explored (amongst others):

- Night time illumination
- Establishment count
- Property prices
- Tele-density

- Vehicle density, etc.

Solution to include:

- Data pipe-line (to extract data from sources)
- Ranking of pin-codes (or more granular geographical construct)
- Geographical representation of solution on any freely available dashboard solution

4. Real-Time Sales Pitch Assistant via Agentic Workflow

Objective

Develop an AI-driven, real-time sales assistant using an agentic workflow to empower life insurance agents during live call center customer conversations.

Problem statement

Life insurance products are complex and agents training is limited, they often struggle to recall specific product details, accurately compare competitor offerings, or handle unexpected objections during live calls. The goal is to build a low-latency, agentic framework that dynamically analyses customer and agent conversation, demographics & segmentation and provides live, context-aware pitch recommendations, objection handling, and compliance guardrails to improve conversion rates and customer experience.

Participants must design a real-time AI pipeline + agentic orchestration model that:

- Ingests live conversation streams (voice-to-text or chat).
- Processes customer intent, sentiment, and context on the fly.
- Orchestrates specialized AI agents (e.g., Product Expert Agent, Objection Handler Agent, Compliance Agent) to generate responses.
- Produces actionable, real-time on-screen prompts and pitch recommendations for the human agent.

Solution to include:

Agentic Workflow Design: A multi-agent framework demonstrating how different AI agents (e.g., Knowledge Retriever, Compliance Checker, Competitor product analyzer, Pitch Generator) collaborate simultaneously without causing delays.

Scoring/Ranking Engine: A logic model that ranks the "next best actions" or recommended products based on the live conversation context.

5. L1 Underwriting Case Summary

Objective

To design and develop an intelligent, GenAI-powered underwriter workbench that automatically ingests case documents, extracts critical data points, and presents a consolidated, interactive summary. The solution aims to reduce the underwriter's cognitive load by proactively identifying and visually flagging risk markers based on standard underwriting guidelines, improving decision-making speed and accuracy.

Problem statement

In the current underwriting process, risk assessment is a highly manual, time-consuming, and fragmented task. An underwriter is required to review a high volume of unstructured and semi-structured documents spanning broadly across three categories:

1. **Application Form Details:** Age, gender, occupation, lifestyle habits, and health declarations.
2. **Financial Details:** Income Tax Returns (ITR), salary slips, Form 16, and bank statements.
3. **Medical Details:** Pathology reports, ECGs, medical examination reports, and attending physician statements.

Apart from the above documents there can also be risk investigation reports, external data source information from credit bureaus, regulatory bodies, reinsurer etc.

Currently underwriters must manually open multiple PDFs and navigate to analyze this data. While the existing frontend displays basic application and health details, it acts only as a static viewer and fails to automatically cross-reference data against underwriting guidelines or visually flag critical risk factors (e.g., classifying an occupation as hazardous, noting an income-to-coverage mismatch, or identifying abnormal medical test levels). This manual synthesis leads to prolonged turnaround times, high context-switching, and an increased risk of human error.

Core Project Scope

- **Data Ingestion & Parsing:** Utilizing GenAI/OCR capabilities to read and extract key entities from sample PDFs and tables across the three primary categories (Application, Financial, Medical).
- **Intelligent Cross-Referencing:** Employing a Large Language Model (LLM) integrated with a knowledge base of sample underwriting guidelines to evaluate the extracted data.

- **Risk Flagging & Reasoning:** Generating automated risk flags (e.g., High, Medium, Low) for specific data points, complete with a GenAI-generated plain-text explanation of *why* the flag was raised based on the guidelines.
- **Interactive Frontend Prototype:** Developing a centralized UI dashboard where underwriters can view the summarized case. The UI will feature visual highlights (e.g., color-coded alerts) for risk markers, clickable elements to see the GenAI reasoning, and side-by-side document verification

Key Deliverable

- **GenAI Extraction & Summary Engine (Backend):** A pipeline (potentially using tools like LangChain, Python, and an LLM API) that takes raw application, financial, and medical PDFs and outputs a structured JSON summary of the applicant.
- **Underwriting Rules/Reasoning Module:** A prompt-engineered or RAG (Retrieval-Augmented Generation) based module that compares the extracted profile against a provided set of underwriting guidelines to identify anomalies or risks.
- **Interactive Underwriter Dashboard (Frontend):** A working UI prototype that displays:
 - A holistic 360-degree view of the applicant.
 - Visual risk indicators (e.g., a red highlight on "Deep Sea Diver" under occupation).
 - Reasoning for explaining the flags

6. Video-Based Lie Detector

Objective

Build a system that automatically determines whether a given video is manipulated or genuine — and if genuine, whether the subject in the video is being deceptive. The system must go beyond single-channel analysis and reason across every aspect of a video: visuals, audio, scene, body language, and metadata. The output must be a clear verdict supported by evidence and explainable reasoning.

Problem Statement

AI-generated and digitally manipulated videos are increasingly difficult to distinguish from real ones. Deepfakes, cloned voices, composite backgrounds, and face-swapped content are now widely accessible to create, yet nearly impossible to detect with the naked eye.

This creates a serious problem in high-stakes situations — insurance claim interviews, identity verification calls, legal proceedings, and investigative contexts — where video is accepted as credible evidence.

Existing detection tools fall short because they are siloed (checking only the face or only the audio, never both together), fail on footage that has been compressed or re-shared, and offer no explanation for their verdict — making them unreliable in any real-world application.

Participants must build a unified system that detects manipulation and deception across all possible video channels simultaneously and explains its reasoning clearly.

Core Project Scope & Deliverables

The solution must cover all four detection areas. Each is a distinct technical challenge — plan accordingly.

Visual Analysis

- Detect face-swap artifacts — unnatural blending, texture inconsistencies at face boundaries
- Identify facial reenactment — expressions or lip movements that do not match the original subject
- Analyse eye blink patterns — frequency and naturalness as an indicator of manipulation
- Detect micro-expressions — brief, involuntary facial movements that are impossible to fake
- Check temporal consistency — flickering, unnatural motion, or frame-level anomalies across the video

Audio & Voice Analysis

- Detect cloned or synthesised voice — unnatural speech rhythm, missing breath sounds, robotic tone
- Verify lip-sync alignment — check whether the audio matches the speaker's mouth movements
- Identify inconsistent background audio — spliced or mismatched ambient sound
- Flag emotional tone mismatch — voice that does not match the nature of what is being said

Scene & Background Integrity

- Detect virtual or composited backgrounds — edge artifacts, inconsistent lighting between subject and scene
- Verify shadow and light source consistency across the frame
- Analyse video metadata — device information, encoding fingerprints, and timestamp inconsistencies

Behavioural Deception Analysis

- Analyse gaze patterns — frequency and direction of eye contact breaks
- Detect body language signals — self-touching, postural shifts, concealed hand movements
- Identify linguistic deception markers in speech — hedging, over-qualification, unnatural phrasing

Deliverables

- A working detection pipeline that accepts a video file and produces a structured verdict
- A confidence output for each detection area — visual, audio, scene, and behavioural
- An explainability layer — showing which frames or regions are suspicious and why
- An API endpoint for integration with external systems
- Evaluation of the system on a publicly available deepfake dataset
- A demonstration video and presentation deck

Bottlenecks

- Videos shared on social media are re-encoded multiple times — manipulation artifacts get destroyed in the process, making detection significantly harder
- A model that works well on controlled datasets often fails on real-world footage — handling this gap is a core challenge
- Combining signals from audio and video is non-trivial — simply merging the outputs of two models does not produce a reliable combined verdict
- Advanced deepfakes are specifically designed to fool detection models — the system must account for this
- Low-quality or low-resolution video is common in real scenarios — the system must remain functional regardless
- Behavioural signals like gaze and expression vary naturally across individuals — the system must not penalise genuine variation

Key Issues to Solve

- How does the system detect manipulation in a video that has been re-shared and re-encoded multiple times?
- How do you tell apart a genuine micro-expression from a visual artifact caused by compression?
- How do you combine signals from four independent detection areas into a single, trustworthy verdict?
- What does the system output when it is not confident — does it communicate uncertainty clearly?
- How do you ensure the behavioural analysis does not unfairly flag individuals with conditions that naturally affect gaze or expression?
- Can the system indicate exactly which part of a video — which frames, which region — has been tampered with?

Impact on Business

- Insurance & Claims: Identify fraudulent video evidence submitted with claims before it influences decisions
- Identity Verification: Detect fake or manipulated videos used in remote onboarding and video-KYC processes

7. Issuance workflow replication using Agentic AI

Objective:

To close the gap between customer expectation and organization issuance process, a new process MinuteAssure is proposed for allowing customers to receive their policy instantly

Details:

Insurance products can be classified in three broad categories:

- Savings/Investment policies
- Term policies
- Annuity policies

While Annuity comes with zero risk and term policies requires medical & financial scrutiny depending on the risk involved, savings policies on the other hand, are often seen in between both. Customers typically perceive them as an investment plan, like MFs & FDs. Today, when a customer wants to purchase a savings policy, their underlying motivation is often to invest money to grow it over time and achieve long-term growth. With the rise of online platforms, the process of purchasing MFs and FDs has become streamlined and digital-first. Customers now expect a similar digital experience when buying life insurance, with easy online applications, paperless documentation, quick policy issuance, and digital tools to manage policies.

After shortlisting the policy and making payment, customers generally want the process to be completed in no time and with minimum effort. Applicants do not want to put effort into arranging and uploading documents. However, organization while issuing, also has a risk involved in terms of sum assured with these policies, which

is typically 10X as of premium paid. From an organization point of view, issuance process should be carried out with due diligence, where insurance organization need to validate KYC, income, contact ability and health parameters etc., with an objective that policies are not issued to fraudulent or unsuitable customers. Given the issuance process followed, this takes time.

To close the gap between customer expectation and organization issuance process, if all these parameters can be validated in real time, real-time issuance experience can be provided, allowing customers to receive their policy instantly.

Approach:

The applicant begins its journey by completing a form that includes basic details of cover wanted, other policy information, personal information, contact details, health questions, and demographic data. If a customer submits their KYC details via consent mechanism such as CKYC or eKYC for identity verification, a background process will be initiated to verify the similarity between the photo provided and photo detected during digital process.

Customer is then prompted to access a camera link for a liveness check (optional for offline, maybe mandatory for online), while simultaneously, we conduct income estimation through any bureau, or there is an option to get it verified from PayU, then provided phone and email information will be verified with a third-party vendor. Finally, the customer's profile risk will be assessed by various risk models to decide if the customer is less risky. Below is the detail of different risk checks which are performed:

1. Identity Fraud Risk: While issuing policy, the applicant must be verified with the government issued ID cards to prevent potential fraud. This verification helps mitigate identity risk and ensures authenticity of the applicant, safeguarding against fraudulent activities during the policy issuance process.

2. Financial Risk: Financial risk is associated with the ability to pay premiums. The applicant's income is estimated to be using a Bureau and verified against the declared income.

3. Medical Risk: The applicant is required to record a video, the same helps in establishing the live proof of the customer. Additionally, a snapshot is captured from the video to compare with the submitted ID documents. This process also serves as a measure against any potential spoofing.

4. Communication Risk: It helps in establishing identity via email or contact number of an applicant. It also identifies any irregularity in contact details, if any.

5. Customer profile Risk:

a.Persistency Risk: PTLL model helps in Identifying the customer profiles who might be at risk of not renewing the policy. This helps in improving persistency.

b.Mortality Risk: We estimate probability of Longevity and general fraud risk linked to the profile by PTLR model at the back end. If the applicant is found risky, Additional due diligence (ADD) may be triggered and applicant follows a detailed issuance process.

c.Customer service Risk: We identify customers using PTCL models, who are more likely to complain within 30 days of issuance of the policy.

d.Free look Risk: PTFM model helps in identifying the customer profiles who might cancel the policy in free look in period. This model helps in reducing free look in cancellation.

Deliverables

- A working pipeline that accepts a sample case and produces insta issuance decision
- Wholebundle of solutions for integration with systems
- Evaluation of the system on a sample of 100 cases
- A demonstration video and presentation deck

8. Bionic Calling

Objective

To design and deploy a real-time **GenAI Co-pilot** that assists insurance agents during live renewal calls by providing sentiment-aware, personalized counterarguments and policy value-propositions, ultimately increasing the 'Retention Ratio' for life insurance premiums while reducing 'Call Handling Time' through instant knowledge retrieval

Problem statement

Insurance agents often face high-stress objections. In a live call, an agent has **less than 3 seconds** to respond before a customer loses interest.

- **Without AI:**
 - The agent relies on memory or generic scripts, often sounding robotic or hesitant.
 - Life insurance policies are complex. A single agent might handle 50 different policy types, each with unique riders and tax benefits.
- **With AI:**
 - GenAI can perform **Real-time Sentiment Analysis**. If it detects "financial stress," it can immediately nudge the agent to offer a "grace period" or a "partial payment" option instead of losing the policy entirely and pops up the exact rebuttal needed
 - AI can instantly pull a specific benefit from a 40-page policy document and translate it into a simple "Why you should care" statement for the customer.
 - AI democratizes expertise. It gives every agent the "playbook" of a 20-year veteran, ensuring a **consistent brand voice** and high persuasion quality across the board

Key components of solutions should be:

1. "Speech-to-text" component to detect real-time customer sentiment & intention
2. Leverage a LLM to generate custom responses based on the above, policy details & customer history
3. Provide the agents with instant pop-ups that suggest the most effective nudge to convince the customer for retention

9. Agentic AI Platform for Claims Processing

Objective

Design and build an autonomous AI platform that manages the complete insurance claims journey — from the moment a claim is raised to the final decision — without any human involvement. The platform must be powered by AI agents that work in coordination, validate each other, and deliver fast, consistent, and auditable outcomes.

Problem Statement

Insurance claims processing today is slow, expensive, inconsistent, and highly dependent on manual effort. Adjusters spend significant time gathering documents, verifying details, interpreting policy terms, and making decisions — all of which introduce delay, human error, and variability.

Despite technology being available, most insurers have only partially automated the process. The real decisions — coverage, fraud, settlement — are still made by humans.

Participants must build a system where AI agents handle every step of the claims lifecycle end-to-end, reliably and explainable, without a human in the loop.

Core Project Scope & Deliverables

The platform must cover the following stages. All stages must be functional and connected — a partial pipeline will not be considered complete.

Intake & Registration

- Accept claims through multiple channels — web form, document upload, or email
- Extract key details: claimant identity, policy reference, incident type, date, and location
- Detect duplicate or repeat submissions automatically
- Send an immediate acknowledgement to the claimant with next steps

Document Verification

- Read and understand submitted documents — medical reports, invoices, FIRs, photographs, repair estimates
- Check for inconsistencies across documents — mismatched dates, names, or amounts
- Flag missing documents and generate follow-up requests automatically

Investigation & Fraud Detection

- Cross-verify claim details against external sources — public records, vehicle data, weather information
- Identify suspicious patterns — repeated claims, unusual timelines, conflicting information
- Generate a fraud risk assessment with clear reasoning, not just a score

Policy Interpretation & Coverage Assessment

- Read and interpret the relevant policy document to determine what is and is not covered
- Apply applicable deductions, limits, and exclusions based on policy terms
- Ensure the assessment is grounded in the actual policy text — not assumptions

Claim Summary & Review

- Produce a structured summary of the claim — incident, documents reviewed, fraud assessment, coverage determination
- Highlight areas of uncertainty clearly so any human reviewer can quickly understand the case

Decision & Settlement

- Issue a clear decision — approved, partially approved, or rejected — with full reasoning
- Generate a decision communication to the claimant in plain language
- Trigger the next action — payment initiation or rejection notice with appeal guidance
- Maintain a complete audit trail of every agent action and decision made

Deliverables

- A working multi-agent system demonstrated on multiple claim scenarios
- An orchestration layer showing how agents communicate and hand off tasks
- A dashboard showing claim status, agent reasoning, and audit log
- End-to-end demonstration: claim submitted to decision issued, fully automated
- Fraud detection capability with supporting evaluation
- Architecture diagram and final presentation

Bottlenecks

- AI models can misinterpret policy language — the system must be grounded strictly in actual policy documents, not general knowledge
- Agents working in parallel can produce conflicting outputs — coordination and conflict resolution must be explicitly designed

- Real-world documents are messy — handwritten forms, low-quality scans, and inconsistent formats must all be handled
- Fraud detection without labelled historical data is difficult — the approach must work with limited supervision
- Every decision must be explainable in plain language — a confidence score alone is not sufficient for regulatory or claimant communication
- The system must handle failures gracefully — if an external data source is unavailable, the pipeline should not break

Key Issues to Solve

- How does the system ensure a downstream agent does not act on an output it should not trust?
- How are contradictions between different data sources or documents resolved?
- How does the system know when to escalate to a human, and how does it hand over context cleanly?
- How is the fraud detection agent able to identify patterns it has not been explicitly trained on?
- How would the system's decision hold up if challenged legally or reviewed by a regulator?
- How do you evaluate whether a natural language decision is correct?

Impact on Business

- Significant reduction in claims processing time — from weeks to hours for straightforward cases
- Lower operational cost by reducing dependence on manual adjuster effort
- More consistent decisions — same facts produce the same outcome every time
- Improved fraud detection through real-time analysis rather than post-payment investigation
- Better claimant experience through faster resolution and clear communication
- Full audit trail for every decision — simplifying compliance and regulatory reporting
- Ability to scale during high-volume periods without additional headcount

10. Intelligent Product Comparison & Gap Analysis Tool

Objective

Develop an automated tool to compare ICICI Prudential (IPru) life insurance products with competitor offerings.

In short: The tool should allow a user to input their desired life insurance features, and the engine will highlight which products from various insurers fit those requirements. This will help the team identify products that meet specific needs, highlight gaps in IPru's current portfolio, and find market "white spaces" for new product opportunities.

Problem statement

Currently, designing new products and benchmarking against competitors is a manual, time-consuming process. It involves reading through long brochures and regulatory guidelines. We need an automated system that uses GenAI to "scan" the market and help us ideate new products by seeing what competitors offer and where IPru can improve.

Project Scope & Key Functionalities

1. Automated Data Management

- **Web Scraping & Extraction:** Build a module to fetch product brochures (PDFs) directly from the official websites of IPru and its peer competitors.
- **Dynamic Refresh:** The tool must automatically refresh the product information every month to ensure the data is always up-to-date.
- **Customer Sentiment (optional):** Integrate customer reviews from the public domain (social media, app stores) to understand which product features are liked or disliked.

2. User-Driven Feature Matching (The "Search" Engine)

- **Input:** The user enters specific life insurance features they are looking for. E.g. *"I am 30 year old and looking for term plan with Return of Premium."*
- **Output:** The tool highlights which products from IPru and different insurers matches the requirements with features.
- **Gap Analysis:** The engine must clearly show the **extent of the gap** between what the user requested and what the available IPru products actually offer.

3. Strategic Peer-to-Peer Comparison

- **Dropdown Selection:** Provide a dropdown menu containing IPru and all peer (competitor) products.
- **Similarity Matching:** If a user selects a specific IPru product and peer product/products, the tool must highlight which products from **other peers** are the closest match to it.
- **Feature Gaps:** For these matches, the tool should display a side-by-side comparison highlighting the gaps between them.

11. Compliance scraper:

Adaptive *Agentic* AI Solution for Automated Regulatory Monitoring

Objective

Create an adaptive Agentic AI engine to track **regulatory updates** from Indian regulatory websites. The goal is to build an automated engine that scrapes new/ updated regulations from the regulatory websites and then send that regulation with its summary via Email to our compliance team.

Problem statement

Currently, team members must manually check these regulatory websites to find new/ updated regulations, which is slow and prone to mistakes.

We need a system which can scrape these regulations and update itself if it encounters any structural changes on the website (Like handling cases when the website is down or website's frontend is updated etc.)

Functional requirements:

- Tracking any updates in the structure of the website.
- Scraping new/ updated regulations if any
- Summarizing the regulations for compliance team
- Updating database capturing all regulations/ notifications/ press releases
- Disseminating the above in a user-friendly manner

Additional capabilities for consideration:

- Delta analysis: If any update is published then it could highlight the updates from the previous version of those regulations.
- **Risk Scoring:** Scoring based on critical indicators such as penalties, mandatory reporting requirements, and capital impact.

Some of the indicative regulators:

IRDAI, PFRDA, BSE, RBI, TRAI etc.

12. Real-time video assistant on website

Today's digital life insurance customer expects a guided, simplified, and conversational journey—especially during complex flows such as buying a new policy. However, ICICI Prudential's current journeys—while robust—face friction arising from dense information, complex disclosures, nuanced health questions, and unclear document requirements.

Objective:

Build an AI-assisted Video-Guided Journey Layer built directly into ICICI Prudential's Buy Onlineflow to:

- Reduce drop-offs and speed up journey completion
- Improve product understanding and customer confidence
- Reduce errors
- Increase digital throughput & adoption by reducing friction
- Create consistent, personalised, interactive experiences

Project Scope & Key Functionalities

- Required to build a webpage similar to ICICI PRUDENTIAL LIFE INSURANCE BOL webpage, by replicating the webpage or mirroring it using mirroring tool.
- Webpage with tooltip which opens a video for customers guidance.
- The video should specifically talk about the relevant field where the customer is stuck

Deliverables:

1. Detailed presentation explaining every aspect of the solution
2. Submission of document with the entire development process
3. Submission of the compiled code used in the project
4. Access to the webpages developed